

From Waves to Insights: A Visual Exploration of Sound

T. Aditya Sai Srinivas^{1*}, Y. Sravanthi², Y. Vinod Kumar³, I.V. Dwaraka Srihith⁴
¹Assistant , Professor,Jayaprakash Narayan College of Engineering, Mahbubnagar
²Student, G Pullaiah College of Engineering & Technology, Kurnool
³ Student, Annamacharya Institute of Technology & Sciences, Rajmpet
⁴ Student, Alliance University, Bangalore

***Corresponding Author**
Email Id: - taditya1033@gmail.com

ABSTRACT

Audio Feature Extraction and Visualization in Python presents a comprehensive framework for uncovering valuable insights from audio data. Leveraging the power of the 'librosa' library, this project demonstrates the extraction of essential audio features, including MFCCs, Chroma, Spectral Contrast, Zero Crossing Rate, and Spectral Centroid. Through insightful visualizations, this work showcases the transformation of raw audio into meaningful data representations. These extracted features offer a deeper understanding of the audio content, enabling applications in speech recognition, music analysis, and beyond. This project invites researchers, developers, and enthusiasts to embark on an auditory journey, revealing the hidden patterns within the realm of sound.

Keywords:- Audio Feature Extraction, Python, Librosa Library, Visualization, MFCCs, Spectral Contrast, Data Analysis.

INTRODUCTION

In the realm of digital audio processing and analysis, the ability to extract meaningful information from raw audio signals is of paramount importance. Whether it's for speech recognition, music classification, or understanding the acoustic characteristics of environmental sounds, the process of audio feature extraction serves as the foundational step in converting complex auditory data into manageable and insightful representations. This journey into the world of "Audio Feature Extraction and Visualization in Python" explores the methodologies and techniques for extracting critical audio features using the versatile 'librosa' library. These features, such as Mel-Frequency Cepstral Coefficients (MFCCs), Chroma, Spectral Contrast, Zero Crossing Rate, and Spectral Centroid, offer a window into the underlying structure and content of audio signals.

Python, with its rich ecosystem of libraries and tools, provides an ideal platform for conducting audio analysis. The 'librosa' library, in particular, is widely adopted for its efficiency and ease of use in audio processing tasks. By delving into this domain, we aim to empower researchers, developers, and enthusiasts to unlock the latent potential within audio data.

This exploration goes beyond the abstract notion of sound and transforms it into a tangible form of data. Through engaging visualizations, we shed light on the intricate patterns and characteristics that exist within audio samples. Each feature extracted contributes to a holistic understanding of the audio's timbral, spectral, and temporal aspects. As we embark on this auditory journey, we invite you to discover the magic of deciphering the language of sound and harnessing it for diverse applications in the world of digital signal processing.

AUDIO FEATURE EXTRACTION

Audio feature extraction is a crucial step in many audio processing tasks, such as speech recognition, music classification, and more. Python provides several libraries that can be used to extract various

audio features from audio samples. In this example, I'll show you how to extract some common audio features using the `librosa` library, which is a popular library for audio analysis. You'll need to install it if you haven't already:

```
pip install librosa
```

```
Requirement already satisfied: librosa in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: audioread>=2.1.9 in /usr/local/lib/python3.10/
Requirement already satisfied: numpy!=1.22.0,!>=1.22.1,!>=1.22.2,>=1.20.3 in /u
Requirement already satisfied: scipy>=1.2.0 in /usr/local/lib/python3.10/dist
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3
Requirement already satisfied: joblib>=0.14 in /usr/local/lib/python3.10/dist
Requirement already satisfied: decorator>=4.3.0 in /usr/local/lib/python3.10/
Requirement already satisfied: numba>=0.51.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: soundfile>=0.12.1 in /usr/local/lib/python3.10
Requirement already satisfied: pooch>=1.0 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: soxr>=0.3.2 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: typing-extensions>=4.1.1 in /usr/local/lib/pyt
Requirement already satisfied: lazy-loader>=0.1 in /usr/local/lib/python3.10/
```

Here's a step-by-step guide to extract audio features from a sample audio file using Python and `librosa`:

```
import librosa
import librosa.display
import matplotlib.pyplot as plt
import numpy as np

# Load an audio file
audio_file = '/content/audio.wav' # Replace with your audio file path
y, sr = librosa.load(audio_file)

# 1. Extract Mel-Frequency Cepstral Coefficients (MFCCs)
mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=13)

# 2. Extract Chroma feature
chroma = librosa.feature.chroma_stft(y=y, sr=sr)

# 3. Extract Spectral Contrast
spectral_contrast = librosa.feature.spectral_contrast(y=y, sr=sr)

# 4. Extract Zero Crossing Rate
zero_crossing_rate = librosa.feature.zero_crossing_rate(y)

# 5. Extract Spectral Centroid
spectral_centroid = librosa.feature.spectral_centroid(y=y, sr=sr)
```

```
# Plot the extracted features
plt.figure(figsize=(12, 8))

# 1. Plot MFCCs
plt.subplot(5, 1, 1)
librosa.display.specshow(mfccs, x_axis='time')
plt.colorbar()
plt.title('MFCCs')

# 2. Plot Chroma
plt.subplot(5, 1, 2)
librosa.display.specshow(chroma, x_axis='time')
plt.colorbar()
plt.title('Chroma')

# 3. Plot Spectral Contrast
plt.subplot(5, 1, 3)
librosa.display.specshow(spectral_contrast, x_axis='time')
plt.colorbar()
plt.title('Spectral Contrast')
```

```
# 4. Plot Zero Crossing Rate
plt.subplot(5, 1, 4)
plt.plot(np.arange(0, len(zero_crossing_rate[0])) * len(y) / len(zero_crossing_rate[0]), zero_crossing_rate[0])
plt.title('Zero Crossing Rate')

# 5. Plot Spectral Centroid
plt.subplot(5, 1, 5)
plt.semilogy(spectral_centroid.T, label='Spectral Centroid')
plt.ylabel('Hz')
plt.xticks([])
plt.xlim([0, spectral_centroid.shape[-1]])
plt.legend()

plt.tight_layout()
plt.show()
```

This code does the following:

- ✓ Loads an audio file using librosa.
- ✓ Extracts several audio features like MFCCs, Chroma, Spectral Contrast, Zero Crossing Rate, and Spectral Centroid using librosa functions.
- ✓ Plots these extracted features for visualization.

MFCCs (Mel-Frequency Cepstral Coefficients) are a widely used and powerful feature extraction technique in audio and speech signal processing. They are especially prevalent in tasks like speech recognition, speaker identification, and music genre classification. MFCCs are

derived from the spectral content of an audio signal and are designed to mimic the human auditory system's perception of sound.

Here's a breakdown of the key concepts related to MFCCs:

1. **Mel-Frequency Scale:** The human auditory system does not perceive all frequencies equally. The Mel scale is a logarithmic scale that models how humans perceive pitch. It is designed to be more in line with human perception than the linear frequency scale.

2. **Filter Banks:** MFCCs are computed by applying a series of overlapping triangular filters on the power spectrum of an audio signal. These filters are spaced according to the Mel scale, with each filter representing a specific frequency range.

3. **Logarithmic Compression:** After applying the filter banks, the logarithm of the filter bank energies is taken. This step further aligns the computation with human perception, as our ears respond to changes in loudness (amplitude) in a logarithmic fashion.

4. **Discrete Cosine Transform (DCT):** The final step involves taking the Discrete Cosine Transform of the log-filter bank energies. This yields a set of coefficients, the MFCCs, which represent the spectral content of the audio signal in a compact and decorrelated form.

MFCCs are typically used as feature vectors for audio analysis tasks. They capture essential information about the shape of the spectral envelope of the audio signal. Usually, only the first few MFCC coefficients (e.g., 13 coefficients) are retained for analysis, as they contain the most discriminative information.

In brief, MFCCs are a valuable tool for representing audio signals in a way that is more informative for certain tasks than raw audio data. They capture both spectral and temporal characteristics of the signal, making them a fundamental feature extraction method in the field of audio and speech processing.

Chroma: Chroma features, often referred to as Chroma vectors or Chromagrams, are a type of audio feature that focus on the pitch content of an audio signal. Chroma features are widely used in music analysis, particularly for tasks related to music genre classification, chord recognition, and melody extraction. They provide a concise representation of the harmonic content of an audio signal by mapping it to the 12 different pitch classes in the Western musical scale.

Here are the key concepts related to Chroma features:

1. **Pitch Class:** In music theory, the Western musical scale consists of 12 pitch classes (C, C#, D, D#, E, F, F#, G, G#, A, A#, B), which correspond to the 12 different notes in an octave.

2. **Chromagram:** A Chromagram is a representation of an audio signal in which each of the 12 pitch classes is assigned a value that represents its presence or strength in the signal. The values are often computed using techniques like Short-Time Fourier Transform (STFT) and then converted into Chroma features.

3. **Chroma Feature Extraction:** To compute Chroma features, the audio signal is first divided into short, overlapping time frames, and for each frame, a Fourier Transform is applied to obtain the spectral information. Next, the spectral information is mapped to the 12 pitch classes using a pitch-class mapping function.

4. **Summation or Normalization:** After mapping the spectral information to pitch classes for each time frame, the values for each pitch class are typically summed or normalized over time frames to produce a single Chroma vector that represents the overall harmonic content of the audio.

Chroma features are particularly useful in music-related tasks because they capture information about the underlying harmony and chord progression of a piece of music. They are often used in conjunction with other features like MFCCs (Mel-Frequency Cepstral Coefficients) to provide a richer description of audio content.

In short, Chroma features are a valuable tool for analyzing the harmonic content of audio signals, especially in music analysis and music information retrieval applications. They provide a compact representation of pitch content, making them useful for tasks involving chord recognition, genre classification, and melody extraction in music.

Spectral Contrast is an audio feature extraction technique used to capture and quantify the differences in amplitude between peaks and valleys in the spectral content of an audio signal. This feature is valuable in various audio processing and analysis tasks, including music genre classification, environmental sound classification, and speech processing. Spectral Contrast provides insights into the texture, timbre, and rhythm of an audio signal.

Here are the key concepts related to Spectral Contrast:

1. **Spectral Analysis:** Spectral Contrast is computed from the spectrum of an audio signal. The spectrum represents how the energy of the audio signal is distributed across different frequency bands.
2. **Frequency Bands:** The spectrum is divided into multiple frequency bands or sub-bands. These bands can be evenly spaced or follow a specific auditory scale (e.g., Mel scale) to better align with human perception.
3. **Calculation of Contrast:** For each frequency band, Spectral Contrast measures the difference in amplitude between the peaks (usually defined as local maxima) and valleys (usually defined as local minima) within that band. This contrast value is computed for each sub-band.
4. **Feature Vector:** After calculating the contrast values for all sub-bands, they are typically organized into a feature vector. The feature vector can be used as input to machine learning algorithms for various audio analysis tasks.
5. **Applications:** Spectral Contrast is particularly useful in tasks where timbral characteristics and texture of audio are important. For example, it can help differentiate between sounds with sharp transients and those with more sustained energy, which can be useful in classifying musical genres or recognizing environmental sounds.

6. **Visualization:** Spectral Contrast can be visualized as a set of values representing the contrast in each frequency band, providing a concise summary of the spectral dynamics of an audio signal.

In short, Spectral Contrast is a valuable audio feature that provides information about the variations in spectral content within different frequency bands of an audio signal. It is often used in audio analysis and classification tasks to capture nuances in the acoustic characteristics of sounds, making it a valuable tool for tasks such as music genre classification and sound event detection.

Zero Crossing Rate and Spectral Centroid are two important audio features used in audio signal processing and analysis. They capture specific characteristics of audio signals that are useful in various applications, including speech recognition, music analysis, and sound classification.

1. Zero Crossing Rate (ZCR):

- **Definition:** ZCR is a measure of how frequently the audio signal changes its sign (from positive to negative or vice versa) within a given time frame. It is a simple yet valuable feature for distinguishing between different types of sounds.

- **Calculation:** To compute ZCR, the audio signal is divided into short, overlapping frames, typically using a technique called frame segmentation. Then, for each frame, the number of times the signal crosses the zero amplitude axis is counted.

- **Applications:** ZCR is useful for differentiating between voiced and unvoiced speech segments, detecting percussive elements in music, and identifying abrupt changes in audio signals, such as silence or noise bursts.

2. Spectral Centroid:

- **Definition:** Spectral Centroid is a feature that provides information about the "center of mass" or the weighted average frequency of an audio signal's spectrum. It

can give insights into the perceived brightness or tonal quality of the sound.


- Calculation: To calculate the Spectral Centroid, the audio signal is first transformed into its spectral representation (e.g., using the Fourier Transform). Then, for each frame or segment of the signal, the Spectral Centroid is computed as the weighted average of the spectral frequencies, with each frequency weighted by its magnitude.

- Applications: Spectral Centroid is useful in music analysis for distinguishing between different musical instruments

(e.g., high-pitched versus low-pitched instruments) and in speech processing for differentiating between voiced and unvoiced speech. It can also be used for content-based audio retrieval and classification.

Both Zero Crossing Rate and Spectral Centroid are fundamental audio features that provide valuable information for a wide range of audio analysis tasks. When combined with other features, they can contribute to more robust and accurate audio analysis and classification systems.

Output:



CONCLUSION

Audio feature extraction is a crucial step in the analysis and understanding of audio signals. We explored two fundamental audio features, Zero Crossing Rate and Spectral Centroid, which offer valuable insights into the temporal and spectral characteristics of audio data. Zero

Crossing Rate helps in identifying transitions and percussive elements, while Spectral Centroid offers information about tonal quality and timbral attributes. These features find applications in speech recognition, music analysis, and sound classification, enhancing the performance of various audio processing tasks. By

leveraging these features alongside others, we can unlock the potential of audio data, enabling the development of more accurate and robust audio-based systems.

REFERENCES

1. <https://towardsdatascience.com/extract-features-of-music-75a3f9bc265d>
2. <https://towardsdatascience.com/get-to-know-audio-feature-extraction-in-python-a499fdaefe42>
3. <https://www.kaggle.com/code/ashishpatel26/feature-extraction-from-audio>
4. <https://maelfabien.github.io/machinelearning/Speech9/>